SemEval 2026 Task 7:

Everyday Knowledge Across Diverse Languages and Cultures

BLEnD

https://github.com/BLEnD-SemEval2026/SemEval-2026-Task-7

Organisers

Nedjma Ousidhoum, Junho Myung, Carla Perez-Almendros, Jiho Jin, Amr Keleg, Meriem Beloucif, Yi Zhou, Rodrigo Agerri, Vladimir Araujo, Naomi Baes, James Barry, Joanne Boisson, Nancy F Chen, Christine de Kock, Aleksandra Edwards, Joseba Fernandez de Landa, Mohamed Fazli Imam, Huda Hakami, Shu-Kai Hsieh, Joseph Marvin Imperial, Roy Ka-Wei Lee, Zhengyuan Liu, Chenyang Lyu, Younes Samih, Johan Sjons, Bryan Tan, Asahi Ushio, Weihua Zheng, Alice Oh, Jose Camacho-Collados

Presenter Nedjma Ousidhoum

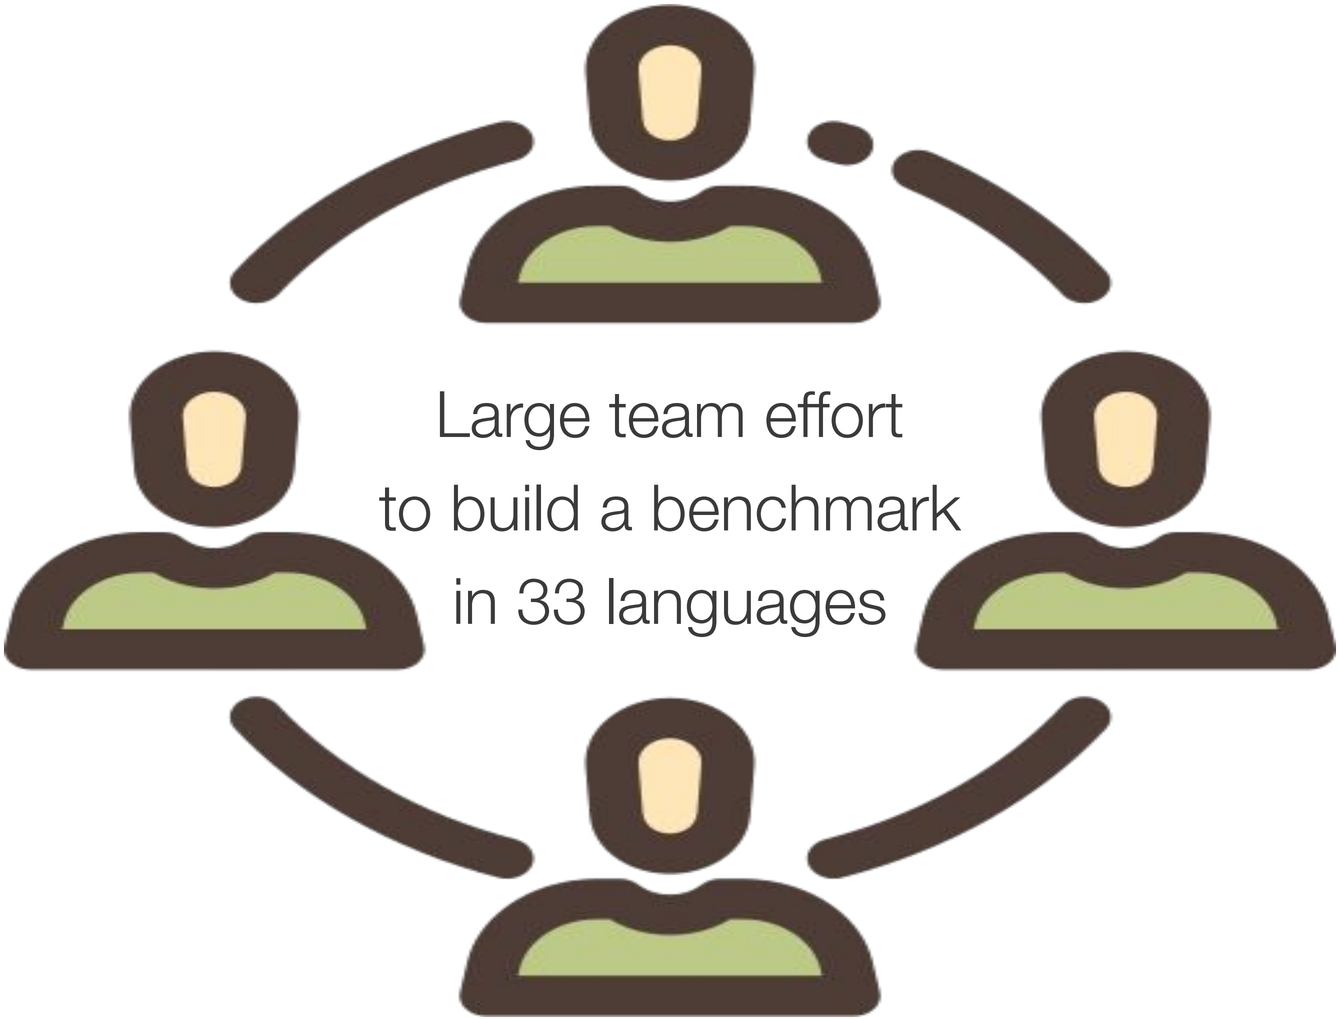
Large team effort
to build a benchmark
in 33 languages

Objective

- Evaluate cultural adaptability across languages
- BLEnD benchmark used exclusively for validation and testing
    - Withheld from training to assess generalisation to unseen, culturally diverse contexts

Data Coverage

- **Original BLEnD (16 culture-language pairs):** Amharic (Ethiopia, am-ET), Arabic (Algeria, ar-DZ), Assamese (Assam, India, as-AS), Azerbaijani (Azerbaijan, az-AZ), Greek (Greece, el-GR), English (United Kingdom, en-GB; USA, en-US), Spanish (Spain, es-ES; Mexico, es-MX), Persian (Iran, fa-IR), Hausa (Northern Nigeria, ha-NG), Indonesian (Indonesia, id-ID), Korean (North Korea, ko-KP; South Korea, ko-KR), Javanese (West Java, Indonesia, su-JB), and Mandarin (China, zh-CN).

Data Coverage

- **Original BLEnD (16 culture-language pairs):** Amharic (Ethiopia, am-ET), Arabic (Algeria, ar-DZ), Assamese (Assam, India, as-AS), Azerbaijani (Azerbaijan, az-AZ), Greek (Greece, el-GR), English (United Kingdom, en-GB; USA, en-US), Spanish (Spain, es-ES; Mexico, es-MX), Persian (Iran, fa-IR), Hausa (Northern Nigeria, ha-NG), Indonesian (Indonesia, id-ID), Korean (North Korea, ko-KP; South Korea, ko-KR), Javanese (West Java, Indonesia, su-JB), and Mandarin (China, zh-CN).

- **For this shared task, we added 17 new pairs:** Arabic (Egypt, ar-EG; Morocco, ar-MA; Saudi Arabia, ar-SA), Bulgarian (Bulgaria, bg-BG), English (Australia, en-AU; Singapore, en-SG), Spanish (Ecuador, es-EC), Basque (Basque Country, eu-PV), French (France, fr-FR), Irish (Ireland, ga-IE), Japanese (Japan, ja-JP), Swedish (Sweden, sv-SE), Tamil (Sri Lanka, ta-LK; Singapore, ta-SG), Tagalog (Philippines, tl-PH), Singaporean Mandarin (Singapore, zh-SG), Taiwanese Mandarin (Taiwan, zh-TW), and Malay (Singapore, ms-SG).

Data Coverage

BLEnD primarily covers

- 33 languages in total
- Low-resource languages from diverse continents
  E.g., Tagalog, Irish, …
- Different language-culture pairs of a language such as:
  - Spanish (EC, ES, MX)
  - English (AU, GB, US)
  - Arabic (DZ, EG, MA, SA)

BLEnD Benchmark

- Data instances consist of question–answer pairs.
- Answers are either:
    - Short free-text responses (Track A), or

Question (in French):
Où est-ce que les étudiants universitaires déjeunent en France ?

Question (Translation):
Where do university students have lunch in France?

Answer:
CROUS

BLEnD

BLEnD Benchmark

- Data instances consist of question–answer pairs.
- Answers are either:
  - Short free-text responses (Track A), or
  - Multiple-choice options (Track B)

Question: What is the most popular sport played in a team at school in Egypt?

BLEnD

√ **A. Football**

X ~~B. baseball~~

X ~~C. field hockey~~

X ~~D. ultimate frisbee~~

Data Creation – Extending the BLEnD Benchmark

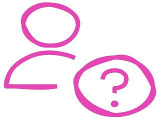

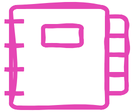

1. Recruited local speakers from different regions to annotate new datasets and provide answers.

Data Creation – Extending the BLEnD Benchmark

1. Recruited local speakers from different regions to annotate new datasets and provide answers.

- Language leads recruited 5+ native speakers per language–culture pair
- Questions cover culturally grounded topics: food, sport, work life, family, education,holidays, celebrations, and leisure
- Answers:
  - Short, concrete, and culturally informed
  - Up to 3 answers per question allowed

Data Creation – Quality Control

2. Quality Control: LLMs and search engines were not allowed; native speakers reviewed the responses.

Data Creation – Quality Control

2. Quality Control: LLMs and search engines were not allowed; native speakers reviewed the responses.

- Native speakers provided culturally grounded answers
- Quality control included qualification tests, attention checks, and demographic verification
- Responses were aggregated by normalising equivalent answers and removing invalid ones
- Answers were translated into English, reviewed by native speakers, and supplemented with synonyms
- Multiple-choice questions were verified to ensure a single correct answer

Final Dataset

Recruited local speakers from different regions to annotate new datasets and provide answers.

Quality Control: LLMs and search engines were not allowed; native speakers reviewed the responses.

Final BLEnD : Extended to 17 additional languages (>30 in total).

Participation Overview

- 140+ registered participants and 880+ submissions
- 22 teams submitted final official systems
- GPT-4.1 and Qwen3-2 in zero-shot settings were used as baselines
- MCQ track received more submissions than SAQ
- Performance varied substantially across languages

Tracks

- Track A (SAQ):

  Generate culturally appropriate short answers

- Track B (MCQ):

  Select the culturally appropriate answer from four options

Track A – Overview

Question (in French):
Où est-ce que les étudiants universitaires déjeunent en France ?

BLEnD

Question (Translation):
Where do university students have lunch in France?

Answer:
CROUS

- English and native-language variants
- Evaluated using exact match with language-specific normalisation
  - Lemmatisation and stemming for highly inflectional languages (e.g. Arabic)
  - Accent removal for languages such as Spanish, French, and Greek
- Semantic similarity not used due to cultural sensitivity of answers
  - Semantically similar responses may still be culturally incorrect

Track B Overview

Question: What is the most popular sport played in a team at school in Egypt?

BLEnD

√ **A. Football**
X ~~B. baseball~~
X ~~C. field hockey~~
X ~~D. ultimate frisbee~~

- Questions are provided in English
- Evaluation is based on answer accuracy
    - All options provided in the annotation may appear in different MCQA pairs
- Accuracy is determined by whether the selected option matches the gold answer
    - Only one valid answer is present among the provided answer options

Track A – Results

king001	59.5
GPT4.1	59.5
qinchihongye	58.3
K-NLPers	55.1
GUIR	54.4

- GPT-4.1 proved to be a strong baseline
  - It matched the top system on average performance
  - It outperformed submitted systems in several languages

Track A – Results

king001	59.5
GPT4.1	59.5
qinchihongye	58.3
K-NLPers	55.1
GUIR	54.4

- GPT-4.1 proved to be a strong baseline
    - It matched the top system on average performance
    - It outperformed submitted systems in several languages

- Some languages were considerably more challenging than others
    - E.g., the best score for Singaporean English was 82.2, 43.2 for Sri Lankan Tamil, compared with 36.4 for Moroccan Arabic

Track A – Takeaways

- Most teams:
  - Used large multilingual LLMs (e.g., Gemini, GPT, Qwen, DeepSeek, and LLaMA)
  - Did not focus on specific languages

Track A – Takeaways

- Most teams:
  - Used large multilingual LLMs (e.g., Gemini, GPT, Qwen, DeepSeek, and LLaMA).
  - Did not focus on specific languages

- Common strategies included:
  - Culture-aware prompt engineering (e.g., prompting the model to act as a local resident)
  - Translation-based or multilingual prompting
  - Constrained output formats and post-processing

Track A – Takeaways

- Most teams:
  - Used large multilingual LLMs (e.g., Gemini, GPT, Qwen, DeepSeek, and LLaMA)
  - Did not focus on specific languages

- Common strategies included:
  - Culture-aware prompt engineering (e.g., prompting the model to act as a local resident)
  - Translation-based or multilingual prompting
  - Constrained output formats and post-processing

- More advanced approaches explored:
  - Ensemble and multi-agent debate frameworks
  - Fine-tuning with synthetic data (e.g., using LoRA)
  - Activation steering and language-vector methods

Track B Results

GUIR	94.81
uir-cis-7	94.28
Pinetree	93.12
GPT 4.1	92.69
K-NLPers	90.55

- Track B's scores were higher than Track A's, as expected.
- GPT-4.1 remained a strong baseline but was generally outperformed by the top-ranked teams

Track B Results

- Most teams:
    - Evaluated multiple LLMs across languages; few focused on specific models or languages
    - Relied on prompting and RAG-based methods
    - Gains mainly from cultural grounding and better prompting, not model changes

Track B Results

- Most teams:
  - Evaluated multiple LLMs across languages; few focused on specific models or languages
  - Relied on prompting and RAG-based methods
  - Gains mainly from cultural grounding and better prompting, not model changes

- Common system designs:
  - Standard instruction prompting
  - Chain-of-thought or light reasoning prompts
  - Structured outputs for A/B/C/D answers
  - Region-aware prompting was most effective

Track B Results

- Most teams:
  - Evaluated multiple LLMs across languages; few focused on specific models or languages
  - Relied on prompting and RAG-based methods
  - Gains mainly from cultural grounding and better prompting, not model changes

- Common system designs:
  - Standard instruction prompting
  - Chain-of-thought or light reasoning prompts
  - Structured outputs for A/B/C/D answers
  - Region-aware prompting was most effective

- Other methods explored:
  - Persona-based prompting
  - Intent-conditioned prompting
  - Data augmentation

Takeaways

- Cultural adaptation varies across methods, languages, and model settings
  - E.g., RAG can help in low-resource contexts but hurt high-resource ones

Takeaways

- Cultural adaptation varies across methods, languages, and model settings
    - E.g., RAG can help in low-resource contexts but hurt high-resource ones

- Evaluation remains an open challenge, i.e.,
    - SAQ evaluation is hard for morphologically rich languages
    - MCQ results can depend on option ordering

Takeaways

- Cultural adaptation varies across methods, languages, and model settings
    - E.g., RAG can help in low-resource contexts but hurt high-resource ones
- Evaluation remains an open challenge, i.e.,
    - SAQ evaluation is hard for morphologically rich languages
    - MCQ results can depend on option ordering
- Misalignment is striking in under-resourced and non-Western contexts, i.e.,
    - overgeneralisation and sensitivity to regional variation (e.g., in Arabic)

Takeaways

- Cultural adaptation varies across methods, languages, and model settings
    - E.g., RAG can help in low-resource contexts but hurt high-resource ones
- Evaluation remains an open challenge, i.e.,
    - SAQ evaluation is hard for morphologically rich languages
    - MCQ results can depend on option ordering
- Misalignment is striking in under-resourced and non-Western contexts, i.e.,
    - overgeneralisation and sensitivity to regional variation (e.g., in Arabic)
- Smaller/efficient models can still perform competitively in specific cases

Thank you!

Any questions?